Multivariate Regression

AECN 896-002

Outline

1. Introduction
2. FWL theorem
3. Samll Sample Property

Multivariate Regression: Introduction

Univariate vs Multivariate Regression Models

Univariate

The most important assumption $E[u|x] = 0$ (zero conditional mean) is almost always violated (unless you data comes from randomized experiments) because all the other variables are sitting in the error term, which can be correlated with x.

Multivariate

More independent variables mean less factors left in the error term, which makes the endogeneity problem less severe

Bi-variate vs. Uni-variate

$$\begin{aligned}\text{Bi-variate} \quad wage =& \beta_0 + \beta_1 educ + \beta_2 exper + u_2 \\ \text{Uni-variate} \quad wage =& \beta_0 + \beta_1 educ + u_1 (= u_2 + \beta_2 exper)\end{aligned}$$

What's different?

- **bi-variate**: able to measure the effect of education on wage, holding experience fixed because experience is modeled explicitly (We say *exper* is controlled for.)

- **uni-variate**: $\hat{\beta}_1$ is biased unless experience is uncorrelated with education because experience was in error term

Another example

The impact of per student spending (`expend`) on standardized test score (`avgscore`) at the high school level

$$\begin{aligned} avgscore =& \beta_0 + \beta_1 expend + u_1 (= u_2 + \beta_2 avginc) \\ avgscore =& \beta_0 + \beta_1 expend + \beta_2 avginc + u_2 \end{aligned}$$

Model with two independent variables

More generally,

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + u$$

- β_0: intercept
- β_1: measure the change in y with respect to x_1, holding other factors fixed
- β_2: measure the change in y with respect to x_1, holding other factors fixed

The Crucial Condition (Assumption) for Unbiasedness of the OLS Estimator

Uni-variate

For $y = \beta_0 + \beta_1 x + u$,

$E[u|x] = 0$

Bi-variate

For $y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + u$,

- Mathematically: $E[u|x_1, x_2] = 0$
- Verbally: for any values of x_1 and x_2, the expected value of the unobservables is zero

Mean independence condition: example

In the following wage model,

$$wage = \beta_0 + \beta_1 educ + \beta_2 exper + u$$

Mean independence condition is

$$E[u|educ, exper] = 0$$

Verbally: this condition would be satisfied if innate ability of students is on average unrelated to education level and experience.

The model with k independent variables

Model

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots + \beta_k x_k + u$$

Mean independence assumption?

β_{OLS} (OLS estimators of βs) is unbiased if,

$$E[u|x_1, x_2, \ldots, x_k] = 0$$

Verbally: this condition would be satisfied if the error term is uncorrelated wtih any of the independent variables, $x_1, x_2, \ldots, x_k$.

Deriving OLS estimators

OLS

Find the combination of βs that minimizes the sum of squared residuals

So,

Denoting the collection of $\hat{\beta}$s as $\hat{\theta}(= \{\hat{\beta}_0, \hat{\beta}_1, \ldots, \hat{\beta}_k\})$,

$$Min_{\theta} \sum_{i=1}^{n} \left[y_i - (\hat{\beta}_0 + \hat{\beta}_1 x_{1,i} + \hat{\beta}_2 x_{2,i} + \cdots + \hat{\beta}_k x_{k,i}) \right]^2$$

Find the FOCs by partially differentiating the objective function (sum of squared residuals) wrt each of $\hat{\theta} (= \{\hat{\beta}_0, \hat{\beta}_1, \ldots, \hat{\beta}_k\})$,

$$\begin{aligned}
\sum_{i=1}^{n} (y_i - (\hat{\beta}_0 + \hat{\beta}_1 x_{1,i} + \hat{\beta}_2 x_{2,i} + \cdots + \beta_k x_{k,i}) &= 0 \quad (\hat{\beta}_0) \\
\sum_{i=1}^{n} x_{i,1} \Big[y_i - (\hat{\beta}_0 + \hat{\beta}_1 x_{1,i} + \hat{\beta}_2 x_{2,i} + \cdots + \beta_k x_{k,i}) \Big] &= 0 \quad (\hat{\beta}_1) \\
\sum_{i=1}^{n} x_{i,2} \Big[y_i - (\hat{\beta}_0 + \hat{\beta}_1 x_{1,i} + \hat{\beta}_2 x_{2,i} + \cdots + \beta_k x_{k,i}) \Big] &= 0 \quad (\hat{\beta}_2) \\
&\vdots \\
\sum_{i=1}^{n} x_{i,k} \Big[y_i - (\hat{\beta}_0 + \hat{\beta}_1 x_{1,i} + \hat{\beta}_2 x_{2,i} + \cdots + \beta_k x_{k,i}) \Big] &= 0 \quad (\hat{\beta}_k)
\end{aligned}$$

Or more succinctly,

$$\begin{aligned}
\sum_{i=1}^{n} \hat{u}_i &= 0 \quad (\hat{\beta}_0) \\
\sum_{i=1}^{n} x_{i,1} \hat{u}_i &= 0 \quad (\hat{\beta}_1) \\
\sum_{i=1}^{n} x_{i,2} \hat{u}_i &= 0 \quad (\hat{\beta}_2) \\
&\vdots \\
\sum_{i=1}^{n} x_{i,k} \hat{u}_i &= 0 \quad (\hat{\beta}_k)
\end{aligned}$$

Implementation of multivariate OLS

R code: Implementation in R

```r
#--- load the fixest package ---#
library(fixest)

#--- generate data ---#
N <- 100 # sample size
x1 <- rnorm(N) # independent variable
x2 <- rnorm(N) # independent variable
u <- rnorm(N) # error
y <- 1 + x1 + x2 + u # dependent variable
data <- data.frame(y = y, x1 = x1, x2 = x2)

#--- OLS ---#
reg <- feols(y ~ x1 + x2, data = data)

#* print the results
reg
```

```
## OLS estimation, Dep. Var.: y
## Observations: 100
## Standard-errors: IID
##             Estimate Std. Error  t value   Pr(>|t|)
## (Intercept) 1.006640   0.096153 10.46911  < 2.2e-16 ***
## x1          0.931657   0.095026  9.80420 3.5494e-16 ***
## x2          1.177465   0.097780 12.04197  < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Presenting regression results

When you are asked to present regression results in assignments or your final paper, use the `msummary()` function from the `modelsummary` package.

Example

```
#* load the package (isntall it if you have no
library(modelsummary)

#* run regression
reg_results <- feols(speed ~ dist, data = cars

#* report regression table
msummary(
  reg_results,
  # keep these options as they are
  stars = TRUE,
  gof_omit = "IC|Log|Adj|F|Pseudo|Within"
)
```

	Model 1
(Intercept)	8.284***
	(0.874)
dist	0.166***
	(0.017)
Num.Obs.	50
R2	0.651
Std. errors	IID
* p < 0.1, ** p < 0.05, *** p < 0.01	

Frisch–Waugh–Lovell Theorem

Consider the following simple model,

$$y_i = \beta_0 + \beta_1 x_{1,i} + \beta_2 x_{2,i} + \beta_3 x_{3,i} + u_i$$

Suppose you are interested in estimating only β_1.

Let's consider the following two methods,

Method 1: Regular OLS

Regress y on x_1, x_2, and x_3 with an intercept to estimate β_0, β_1, β_2, β_3 at the same time (just like you normally do)

Method 2: 3-step

- regress y on x_2 and x_3 with an intercept and get residuals, which we call $\hat{u}_y$
- regress x_1 on x_2 and x_3 with an intercept and get residuals, which we call $\hat{u}_{x_1}$
- regress $\hat{u}_y$ on $\hat{u}_{x_1}$ ($\hat{u}_y = \alpha_1 \hat{u}_{x_1} + v_3$)

Frisch-Waugh–Lovell theorem

Methods 1 and 2 produces the same coefficient estimate on x_1

$$\hat{\beta}_1 = \hat{\alpha_1}$$

Partialing out Interpretation from Method 2

Step 1

Regress y on x_2 and x_3 with an intercept and get residuals, which we call $\hat{u}_y$

- $\hat{u}_y$ is void of the impact of x_2 and x_3 on y

Step 2

Regress x_1 on x_2 and x_3 with an intercept and get residuals, which we call $\hat{u}_{x_1}$

- $\hat{u}_{x_1}$ is void of the impact of x_2 and x_3 on x_1

Step 3

Regress $\hat{u}_y$ on $\hat{u}_{x_1}$, which produces an estimte of β_1 that is identical to that you can get from regressin y on x_1, x_2, and x_3

Interpretation

- Regressing y on all explanatory variables (x_1, x_2, and x_3) in a multivariate regression is as if you are looking at the impact of a single explanatory variable with the effects of all the other effects partiled out

- In other words, including variables beyond your variable of interest lets you control for (remove the effect of) other variables, avoiding confusing the impact of the variable of interest with the impact of other variables.

Small Sample Properties of OLS Estimators

Unbiasedness of OLS Estimator

OLS estimators of multivariate models are unbiased under certain conditions

Condition 1

Your model is correct (Assumption $MLR.1$)

Condition 2

Random sampling (Assumption $MLR.2$)

Conditions 3

No perfect collinearity (Assumption $MLR.3$)

Perfect Collinearity

No Perfect Collinearity

Any variable cannot be a linear function of the other variables

Example (silly)

$$wage = \beta_0 + \beta_1 educ + \beta_2(3 \times educ) + u$$

(More on this later when we talk about dummy variables)

Zero Conditional Mean

$$E[u|x_1, x_2, \ldots, x_k] = 0 \quad \text{(Assumption MLR.4)}$$

Unbiasedness of OLS estimators

If all the conditions $MLR.1 \sim MLR.4$ are satisfied, OLS estimators are unbiased.

$$E[\hat{\beta}_j] = \beta_j \quad {}^{\forall} j = 0, 1, \ldots, k$$

Endogeneity (Definition)

$$E[u|x_1, x_2, \ldots, x_k] = f(x_1, x_2, \ldots, x_k) \neq 0$$

What could cause endogeneity problem?

- functional form misspecification

$$\begin{aligned} wage =& \beta_0 + \beta_1 log(x_1) + \beta_2 x_2 + u_1 \quad (\text{true}) \\ wage =& \beta_0 + \beta_1 x_1 + \beta_2 x_2 + u_2 (= log(x_1) - x_1) \quad (\text{yours}) \end{aligned}$$

- omission of variables that are correlated with any of $x_1, x_2, \ldots, x_k$ (more on this soon)
- other sources of enfogeneity later

Variance of the OLS estimators

Homoeskedasticity

$$Var(u|x_1, \ldots, x_k) = \sigma^2 \quad \text{(Assumption MLR.5)}$$

Variance of the OLS estimator

Under conditions $MLR.1$ through $MLR.5$, conditional on the sample values of the independent variables,

$$Var(\hat{\beta}_j) = \frac{\sigma^2}{SST_j(1 - R_j^2)},$$

where $SST_j = \sum_{i=1}^{n}(x_{ji} - \bar{x}_j)^2$ and R_j^2 is the R-squared from regressing x_j on all other independent variables including an intercept. (We will revisit this equation)

Estimating σ^2

Just like uni-variate regression, you need to estimate σ^2 if you want to estimate the variance (and standard deviation) of the OLS estimators.

uni-variate regression

$$\hat{\sigma}^2 = \sum_{i=1}^{N} \frac{\hat{u}_i^2}{n-2}$$

multi-variate regression

A model with k independent variables with intercept.

$$\hat{\sigma}^2 = \sum_{i=1}^{N} \frac{\hat{u}_i^2}{n-(k+1)}$$

You solved $k+1$ simultaneous equations to get $\hat{\beta}_j$ $(j = 0, \ldots, k)$. So, once you know the value of $n-k-1$ of the residuals, you know the rest.

The estimator of the variance of the OLS estimator is therefore

$$\widehat{Var\hat{\beta}}_j = \frac{\hat{\sigma}^2}{SST_j(1-R_j^2)}$$