

# Monte Carlo Simulation

AECN 896-002

# Outline

1. Introduction
2. MC Simulations

# Monte Carlo Simulation: Introduction

---

## Monte Carlo Simulation

A way to test econometric theories via simulation

## How is it used in econometrics?

- confirm econometric theory numerically
  - OLS estimators are unbiased if  $E[u|x] = 0$  along with other conditions (theory)
  - I know the above theory is right, but let's check if it is true numerically
- You kind of sense that something in your data may cause problems, but there is no proven econometric theory about what's gonna happen (I used MC simulation for this purpose a lot)
- assist students in understanding econometric theories by providing actual numbers instead of a series of Greek letters

### Question

Suppose you are interested in checking what happens to OLS estimators if  $E[u|x] = 0$  (the error term and  $x$  are not correlated) is violated.

Can you use the real data to do this?

## Key part of MC simulation

You generate data (you have control over how data are generated)

- You know the true parameter unlike the real data generating process
- You can change only the part that you want to change about data generating process and econometric methods with everything else fixed

# Generating data

## Pseudo random number generators

Algorithms for generating a sequence of numbers whose properties **approximate** the properties of sequences of random numbers




## Examples in R: Uniform Distribution

```
runif(5) # default is min=0 and max=1
```

```
## [1] 0.7276097 0.5903365 0.6655839 0.9031736 0.5203449
```

```
x <- runif(10000)  
hist(x)
```



## Pseudo random number generator

- Pseudo random number generators are not really random number generators
- What numbers you will get are pre-determined
- What numbers you will get can be determined by setting a **seed**

## An example

```
set.seed(2387438)  
runif(5)
```

```
## [1] 0.0474233 0.7116970 0.4066674 0.2422949 0.3567480
```

## Question


What benefits does setting a seed have?

## Examples in R: Normal Distribution

$$x \sim N(0, 1)$$

```
# default is mean = 0, sd = 1  
x <- rnorm(10000)  
hist(x)
```


Histogram of x



$$x \sim N(2, 2)$$

```
# mean = 2, sd = 2  
x <- rnorm(10000, mean = 2, sd = 2)  
hist(x)
```

Histogram of x



## Other distributions

- Beta
- Chi-square
- F
- Logistic
- Log-normal
- many others

**d, p, q, r**


For each distribution, you have four different kinds of functions:

- **d norm**: density function
- **p norm**: distribution function
- **q norm**: quantile function
- **r norm**: random draw

**dnorm**

`dnorm(x)` gives you the height of the density function at  $x$ .

`dnorm(-1)` and `dnorm(2)`







**pnorm**

`pnorm(x)` gives you the probability that a single random draw is **less** than  $x$ .

`pnorm(-1)`



`pnorm(2)`




## Practice

What is the probability that a single random draw from a Normal distribution with `mean = 1` and `sd = 2` is less than 1?

`qnorm(x)`, where  $0 < x < 1$ , gives you a number  $\pi$ , where the probability of observing a number from a single random draw is less than  $\pi$  with probability of  $x$ .

We call the output of `qnorm(x)`,  $x$  quantile of the standard Normal distribution (because the default is `mean = 0` and `sd = 1` for `rnorm()`).

`qnorm(0.95)`



## Practice

What is the 88% quantile of Normal distribution with `mean = 0` and `sd = 9`?

# Monte Carlo Simulation: Introduction

---



## Monte Carlo Simulation: Steps

- specify the data generating process
- generate data based on the data generating process
- get an estimate based on the generated data (e.g. OLS, mean)
- repeat the above steps many many times
- compare your estimates with the true parameter

## Question

Why do the steps 1 – 3 many many times?

# Monte Carlo Simulation: Example 1

Is sample mean really an unbiased estimator of the expected value?

That is, is  $E[\frac{1}{n} \sum_{i=1}^n x_i] = E[x]$ , where  $x_i$  is an independent random draw from the same distribution,

## Sample Mean: Steps 1-3

```
#--- steps 1 and 2: ---#  
# specify the data generating process and generate data  
x <- runif(100) # Here,  $E[x]=0.5$ 
```

```
#--- step 3 ---#  
# calculate sample mean  
mean_x <- mean(x)  
mean_x
```

```
## [1] 0.507078
```

### Sample Mean: Step 4

- repeat the above steps many times
- We use a **loop** to do the same (similar) thing over and over again

## Loop: for loop

```
#--- the number of iterations ---#  
B <- 1000  
  
#--- repeat steps 1-3 B times ---#  
for (i in 1:B) {  
  print(i) # print i  
}
```

## Verbally

For each of  $i$  in  $1 : B$  ( $1, 2, \dots, 1000$ ), do `print(i)`.

- `i` takes the value of 1, and then `print(1)`
- `i` takes the value of 2, and then `print(2)`
- ...
- `i` takes the value of 999, and then `print(999)`
- `i` takes the value of 1000, and then `print(1000)`

## Step 4

```
#--- the number of iterations ---#
B <- 1000

#--- create a storage that stores estimates ---#
estimate_storage_mean <- rep(0, B)

#--- repeat steps 1-3 B times ---#
for (i in 1:B) {
  #--- steps 1 and 2: ---#
  # specify the data generating process and generate data
  x <- runif(100) # Here, E[x]=0.5

  #--- step 3 ---#
  # calculate sample mean
  mean_x <- mean(x)
  estimate_storage_mean[i] <- mean_x
}
```

## Compare your estimates with the true parameter

```
mean(estimate_storage_mean)
```

```
## [1] 0.500199
```

```
hist(estimate_storage_mean)
```

# Monte Carlo Simulation: Example 2

## Question

What happens to  $\beta_1$  if  $E[u|x] \neq 0$  when estimating  $y = \beta_0 + \beta_1 x + u$ ?



```

#--- load the fixest package for feols() ---#
library(fixest)

#--- Preparation ---#
B <- 1000 # the number of iterations
N <- 100 # sample size
estimate_storage <- rep(0, B) # estimates storage


#--- repeat steps 1-3 B times ---#
for (i in 1:B) {
  #--- steps 1 and 2: ---#
  mu <- rnorm(N) # the common term shared by both x and u
  x <- rnorm(N) + mu # independent variable
  u <- rnorm(N) + mu # error
  y <- 1 + x + u # dependent variable
  data <- data.frame(y = y, x = x)

  #--- OLS ---#
  reg <- feols(y ~ x, data = data) # OLS
  estimate_storage[i] <- reg$coefficient["x"]
}

```

```
hist(estimate_storage)
```

## Histogram of estimate\_storage



# Example 3: Variance of OLS Estimators

## Model

$$y = \beta_0 + \beta_1 x + u$$

- $x \sim N(0, 1)$
- $u \sim N(0, 1)$
- $E[u|x] = 0$

## Variance of the OLS estimator

True Variance of  $\hat{\beta}_1$ :  $V(\hat{\beta}_1) = \frac{\sigma^2}{\sum_{i=1}^n (x_i - \bar{x})^2} = \frac{\sigma^2}{SST_X}$

Its estimator:  $\widehat{V(\hat{\beta}_1)} = \frac{\hat{\sigma}^2}{SST_X} = \frac{\sum_{i=1}^n \hat{u}_i^2}{n-2} \times \frac{1}{SST_X}$

## Question

Does the estimator really work? (Is it unbiased?)

```

set.seed(903478)

#--- Preparation ---#
B <- 10000 # the number of iterations
N <- 100 # sample size
beta_storage <- rep(0, B) # estimates storage for beta
V_beta_storage <- rep(0, B) # estimates storage for V(beta)
x <- rnorm(N) # x values are the same for every iteration
SST_X <- sum((x - mean(x))^2)

#--- repeat steps 1-3 B times ---#
for (i in 1:B) {
  #--- steps 1 and 2: ---#
  u <- 2 * rnorm(N) # error
  y <- 1 + x + u # dependent variable
  data <- data.frame(y = y, x = x)

  #--- OLS ---#
  reg <- feols(y ~ x, data = data) # OLS
  beta_storage[i] <- reg$coefficient["x"]
  #* store estimated variance of beta_1_hat
  V_beta_storage[i] <- vcov(reg)["x", "x"]
}

```

## True Variance

- $SST_X = 112.07$
- $\sigma^2 = 4$

$$V(\hat{\beta}) = 4/112.07 = 0.0357$$


## Check

Your Estimates of Variance of  $\hat{\beta}_1$ ?

```
# === mean ===#  
mean(V_beta_storage)
```

```
## [1] 0.03579118
```

```
ggplot(data = data.frame(x = V_beta_storage)) +  
  geom_density(aes(x = x)) +  
  geom_vline(xintercept = round(4 / SST_X, digits = 4))
```



# Exercise

## Problem

Using MC simulations, find out how the variation in  $x$  affects the OLS estimators

## Model setup

$$y = \beta_0 + \beta_1 x_1 + u$$

$$y = \beta_0 + \beta_1 x_2 + u$$

- $x_1 \sim N(0, 1)$  and  $x_2 \sim N(0, 9)$
- $u \sim N(0, 1)$
- $E[u_1|x] = 0$  and  $E[u_2|x] = 0$



# Solution

```
#--- Preparation ---#
B <- 1000 # the number of iterations
N <- 100 # sample size
estimate_storage <- matrix(0, B, 2) # estimates storage

for (i in 1:B) {
  #--- generate data ---#
  x_1 <- rnorm(N, sd = 1) # indep var 1
  x_2 <- rnorm(N, sd = 3) # indep var 2
  u <- rnorm(N) # error
  y_1 <- 1 + x_1 + u # dependent variable 1
  y_2 <- 1 + x_2 + u # dependent variable 2
  data <- data.table(y_1 = y_1, y_2 = y_2, x_1 = x_1, x_2 = x_2)

  #--- OLS ---#
  reg_1 <- feols(y_1 ~ x_1, data = data) # OLS
  reg_2 <- feols(y_2 ~ x_2, data = data) # OLS

  #--- store coef estimates ---#
  estimate_storage[i, 1] <- reg_1$coefficient["x_1"] # equation 1
  estimate_storage[i, 2] <- reg_2$coefficient["x_2"] # equation 2
}
```

```
#--- assign new names ---#
beta_1s <- estimate_storage[, 1]
beta_2s <- estimate_storage[, 2]

#--- mean ---#
mean(beta_1s)
```

# Visualization

```
plot_data_1 <- data.table(x = beta_1s, type = "Equation 1")
plot_data_2 <- data.table(x = beta_2s, type = "Equation 2")
plot_data <- rbind(plot_data_1, plot_data_2)
ggplot(data = plot_data) +
  geom_density(aes(x = x, fill = type), alpha = 0.5) +
  scale_fill_discrete(name = "") +
  xlab("Coefficient Estimate") +
  theme(
    legend.position = "bottom"
  )
```